

Programmazione Avanzata e Paradigmi

Ingegneria e Scienze Informatiche - UNIBO

a.a 2014/2015

Lecturer: Alessandro Ricci

[module lab 2.2] THREAD SAFETY & LIVENESS

THREAD SAFETY DEFINITION

- A central aspect of concurrent programming is writing **thread-safe** code / thread-safe classes
 - a class is *thread-safe* if *it continues to behave correctly when accessed from multiple threads*
 - **regardless of the scheduling or interleaving of the execution of those threads by the runtime environment**
 - **with no additional synchronization or other coordination on the part of the calling code**
- *Correctness* means that a class *conforms to its specification*
 - a good specification defines
 - **invariants** constraining an object's state
 - **post-conditions** describing the effects of its operations
- Thread-safe classes encapsulate any needed synchronization so that clients need not to provide their own

SHARED MUTABLE STATE

- Writing thread-safe code is – at its core – about *managing access to state and in particular to **shared, mutable state***:
 - **shared**: variable or object could be accessed by multiple threads
 - **mutable**: its value could change during its lifetime
- if multiple threads access the same mutable state variable without appropriate synchronization, the program is *broken*
 - race conditions
- There are three ways to fix it:
 - don't share state variable across threads
 - make the state variable immutable
 - use synchronization whenever accessing the state variable

STATELESS OBJECT

- Stateless objects are *always* thread-safe
 - actions of a thread accessing a stateless object cannot affect the correctness of operations in other threads
- Example - a *factorizer service*
 - source: `pap.lab05.factorizer.FactorizerService`

```
public interface FactorizerService {  
    int[] getFactors(long n);  
}
```

- **Stateless** implementation
 - check source:
`pap.lab05.factorizer.StatelessFactorizer`
 - this class is thread-safe *by construction*

STATE-FULL OBJECTS

- Thread safety is undermined as soon as we share and access in R/W stateful objects
 - **mutable** state-full objects, in particular
- Example: adding a *cache* to the factorizer service
 - check source:
`pap.lab03.factorizer.FactorizerWithCache_unsafe`
 - this class is **not** thread-safe:
 - **check & act** problem => race conditions

RACE CONDITIONS

- The concurrent execution of non-atomic sequence of statements that should be considered atomic generate **race conditions**
 - occur when the *correctness of a computation depends on the relative timing or interleaving of multiple threads* by the runtime (...and getting the right answer relies on lucky timing..)
- Main examples
 - **lost updates**
 - when executing concurrently non-atomic *read-modify-write* operations
 - ex: count++
 - **check-and-act**
 - when a potentially state observation is used to make a decision on what to do next
 - example:

| |
|--|
| <pre>If (file X doesn't exist) -- check then create file X -- act</pre> |
|--|
 - since check+act are not atomic, *the state can change after check and before act.*

COMPOUND ATOMIC ACTIONS

- *check-and-act* and *read-modify-write* are examples of compound actions
 - sequences of operations that must be executed atomically in order to remain thread-safe


ATOMIC COMPOUND ACTIONS IN JAVA : SYNCHRONIZED BLOCKS

- Compound-actions - and atomic statement blocks - in Java can be realized by means of synchronized blocks or methods



```
synchronized( lock ) {  
    statement  
    statement  
    statement  
}
```



- A synchronized block has 2 parts 
 - a reference to an object that will serve as the lock
 - block of code to be guarded by the lock
- Mostly used at a *method* level
 - synchronized attribute
 - more about this in next modules when discussing monitors in Java

INTRINSIC LOCK AND ENTRY SET

- Atomic blocks work by exploiting the **lock** embedded in each Java object (more on this in next modules)
 - called **intrinsic lock** or monitor lock
 - functioning as a *guard* for the block
- The lock is automatically acquired and then released by a thread respectively when entering and exiting the block
 - if the lock is already acquired, the thread is blocked (suspended) and added to the entry set
 - when a thread exited the block, one thread of the entry set is selected and re-activated
 - no ordering policy is specified
 - if the lock is not released by the thread inside the block, threads in the entry set are blocked forever (starvation)
- For static methods and fields, the lock is associated to the related Class object
- For **synchronized** methods, the object serving as lock is **this**


LOCK REENTRANCY (1/2)

- def: **lock reentrancy**
 - when locks are acquired on a *per-thread* basis
 - vs. per-invocation basis
 - per-invocation basis is adopted instead as default locking behavior for Pthreads (POSIX threads) mutex-es
- Java intrinsic locks are reentrant:
 - if a thread tries to acquire a lock that it already holds, the request succeeds

LOCK REENTRANCY (2/2)

- Reentrancy facilitates encapsulation of locking behaviour and thus simplify the development of OO concurrent code

```
public class Widget {  
    public synchronized doSomething(){...}  
}  
  
public class LoggingWidget extends Widget {  
    public synchronized void doSomething(){  
        System.out.println(toString()+" : calling doSomething");  
        super.doSomething();  
    }  
}
```



- Without reentrancy the above example would lead to a deadlock

PERFORMANCE: POOR CONCURRENCY PROBLEM

- The misuse of atomic blocks can lead to performance problems.
- Example: over constrained factorizer service
 - check source:
`pap.lab05.factorizer.FactorizerWithCache_overconstrained`
- This solution is thread-safe but not acceptable
 - it enforces a sequentialization of computations that can be done concurrently
 - poor performances
- Careful choice of what parts must be designed and implemented as critical sections
 - examples with the factorizer service with cache
 - `pap.lab05.factorizer.FactorizerWithCache_quite_good`
 - `pap.lab05.factorizer.FactorizerWithCache_good`

AVOIDING LIVENESS HAZARD

- Tension between *safety* and *liveness*
 - using locking to ensure thread safety
 - ...but indiscriminate use of locking can cause "lock-ordering" deadlocks
 - using thread pools and semaphores to bound resource consumption
 - ...but failure to understand activities being bounded can cause resource deadlocks

DEADLOCKS

- A situation wherein *two or more competing actions are waiting for the other to finish, and thus neither ever does*
- Coffman *necessary conditions* for a deadlock to occur (1971)
 - **mutual exclusion condition**
 - a resource that cannot be used by more than one process at a time
 - **hold and wait condition**
 - processes already holding resources may request new resources
 - **no preemption condition**
 - no resource can be removed from a process holding it
 - resources can be released only by the explicit action of the process
 - **circular wait condition**
 - two or more processes form a circular chain where each process waits for a resource that the next process in the chain holds
- Deadlock can only occur in systems where all 4 conditions hold true

DEADLOCKS WITH LOCKS

- It happens when
 - multiple threads wait forever due to cyclic locking dependency
 - simplest case
 - when thread A holds lock L and tries to acquire lock M, but at the same time thread B holds M and tries to acquire L, both thread will wait for ever
 - deadly embrace

DEADLOCKS DETECTION & RECOVERY

- Deadlocks **detection** and **recovery**
 - adopted in databases
 - databases are designed to detect and recover from deadlocks
 - transactions typically acquire many locks, until they commit
 - not so uncommon for two transactions to deadlock
 - identifying the set of transactions that are deadlocked by analyzing *is-waiting* dependency graph
 - looking for cycles
 - if a cycle is found, a victim is selected and the transaction aborted
- No automated deadlock detection / recovery mechanism in JVM
 - if threads deadlock, *that's all folks*
 - we can just shutdown the application
 - “post-mortem” diagnosis support

DEADLOCK DIAGNOSING

- *Thread dump* support provided by the JVM
 - triggered by
 - sending the JVM process a SIGQUIT signal on UNIX (kill -3)
 - pressing CTRL-\ on UNIX
 - pressing CTRL-Break on Windows
- Thread dump content
 - stack trace for each running thread
 - locking information
 - which locks are held by each thread, in which stack frame they were acquired, and which lock a blocked thread is waiting for

LOCK-ORDERING DEADLOCKS


- Simple example:

```
public class LeftRightDeadlock {
    private final Object left = new Object();
    private final Object right = new Object();

    public void leftRight(){
        synchronized(left){
            synchronized(right){
                doSomething();
            }
        }
    }

    public void rightLeft(){
        synchronized(right){
            synchronized(left){
                doSomethingElse();
            }
        }
    }

    private void doSomething(){ System.out.println("something.");}
    private void doSomethingElse(){ System.out.println("something else.");}
}
```



DYNAMIC LOCK-ORDER DEADLOCKS

- When locks to lock are established dynamically
 - basic example: Transfer money between bank accounts

```
public class Test1a {  
  
    private static final int NUM_THREADS = 20;  
    private static final int NUM_ACCOUNTS = 5;  
    private static final int NUM_ITERATIONS = 10000000;  
    private static final Random gen = new Random();  
    private static final Account[] accounts = new Account[NUM_ACCOUNTS];  
  
    public static void transferMoney(Account from, Account to, int amount)  
        throws InsufficientBalanceException {  
        synchronized (from) {  
            synchronized (to) {  
                if (from.getBalance() < amount)  
                    throw new InsufficientBalanceException();  
                from.debit(amount);  
                to.credit(amount);  
            }  
        }  
    }  
  
    ...  
}
```

```

public class Test1a {

    private static final int NUM_THREADS = 20;
    private static final int NUM_ACCOUNTS = 5;
    private static final int NUM_ITERATIONS = 10000000;
    private static final Random gen = new Random();
    private static final Account[] accounts = new Account[NUM_ACCOUNTS];

    public static void transferMoney(Account from, Account to, int amount)
        throws InsufficientBalanceException {...}

    static class TransferThread extends Thread {
        public void run() {
            for (int i = 0; i < NUM_ITERATIONS; i++){
                int fromAcc = gen.nextInt(NUM_ACCOUNTS);
                int toAcc = gen.nextInt(NUM_ACCOUNTS);
                int amount = gen.nextInt(10);
                try {
                    transferMoney(accounts[fromAcc], accounts[toAcc], amount);
                } catch (InsufficientBalanceException ex){}
            }
        }
    }

    public static void main(String[] args) {
        for (int i = 0; i < accounts.length; i++){
            accounts[i] = new Account(1000);
        }
        for (int i = 0; i < NUM_THREADS; i++){
            new TransferThread().start();
        }
    }
}

```

ORDERING LOCKS

- Deadlock came because the two threads attempted to acquire the locks *in a different order*
 - if they asked for the locks in the same order, *there would be no cyclic locking dependency* and therefore no deadlock
- A program will be free of lock-ordering deadlocks *if all threads acquire the locks they need in a fixed global order*
 - verifying consistent lock ordering requires a global analysis of your program's locking behavior

```

public class AccountManager {

    private final Account[] accounts;

    public AccountManager(int nAccounts, int amount){
        accounts = new Account[nAccounts];
        for (int i = 0; i < accounts.length; i++){
            accounts[i] = new Account(amount);
        }
    }

    public void transferMoney(int from, int to, int amount)
                                throws InsufficientBalanceException {

        int first = from;
        int last = to;

        if (first > last){
            last = first;
            first = to;
        }

        synchronized (accounts[first]) {
            synchronized (accounts[last]) {
                if (accounts[from].getBalance() < amount)
                    throw new InsufficientBalanceException();
                accounts[from].debit(amount);
                accounts[to].credit(amount);
            }
        }
    }
}

```

```

public class Test1b {

    private static final int NUM_THREADS = 20;
    private static final int NUM_ACCOUNTS = 5;
    private static final int NUM_ITERATIONS = 10000000;
    private static final Random gen = new Random();

    static class TransferThread extends Thread {
        AccountManager man;
        TransferThread(AccountManager man){
            this.man = man;
        }
        public void run() {
            for (int i = 0; i < NUM_ITERATIONS; i++){
                int fromAcc = gen.nextInt(NUM_ACCOUNTS);
                int toAcc = gen.nextInt(NUM_ACCOUNTS);
                int amount = gen.nextInt(10);
                try {
                    man.transferMoney(fromAcc,toAcc,amount);
                } catch (InsufficientBalanceException ex){
                }
            }
        }
    }

    public static void main(String[] args) {
        AccountManager man = new AccountManager(NUM_ACCOUNTS,1000);
        for (int i = 0; i < NUM_THREADS; i++){
            new TransferThread(man).start();
        }
    }
}

```

DEADLOCKS BETWEEN COOPERATING OBJECTS

- More subtle deadlocks can happen in cooperating objects, in which no methods explicitly acquire two locks, but where this happens indirectly
- A common example: *Observer pattern*
 - observers observing observed object
 - different control flows executing methods for observers and observed
- The more general problem
 - event-oriented pattern implementation in OO languages
 - coupling controls among sources and observers of events
 - MVC case study


```

interface IObserved {
    int getState();
    void register(IObserver obj);
}

class MyEntityA implements IObserved {

    private List<IObserver> obsList;
    private int state;

    public MyEntityA(){
        obsList = new ArrayList<IObserver>();
    }

    public void register(IObserver obs) {
        obsList.add(obs);
    }

    public synchronized int getState() {
        return state;
    }

    public synchronized void changeState1() {
        state++;
        for (IObserver o: obsList){
            o.notifyStateChanged(this);
        }
    }

    public synchronized void changeState2() {
        state--;
        for (IObserver o: obsList){
            o.notifyStateChanged(this);
        }
    }
}

```

```

interface IObserver {
    void notifyStateChanged(IObserved obs);
}

class MyEntityB implements IObserver {

    List<IObserved> obsList;

    public MyEntityB(){
        obsList = new ArrayList<IObserved>();
    }

    public synchronized void observe(IObserved obj){
        obsList.add(obj);
        obj.register(this);
    }

    public synchronized
        void notifyStateChanged(IObserved obs) {
        synchronized(System.out){
            System.out.println(
                "state changed: "+obs.getState());
        }
    }

    public synchronized int getOverallState() {
        int sum = 0;
        for (IObserved o: obsList){
            sum += o.getState();
        }
        return sum;
    }
}

```

```

class MyThreadA extends Thread {
    MyEntityA obj;

    public MyThreadA(MyEntityA obj){
        this.obj = obj;
    }

    public void run(){
        while (true){
            obj.changeState1();
            obj.changeState2();
        }
    }
}

class MyThreadB extends Thread {
    MyEntityB obj;

    public MyThreadB(MyEntityB obj){
        this.obj = obj;
    }

    public void run(){
        while (true){
            log("overall state: "+
                obj.getOverallState());
        }
    }

    private void log(String msg){
        synchronized(System.out){
            System.out.println("[ "+this+" ] "+msg);
        }
    }
}

```

```

public class Test2 {
    public static void main(String[] args) {

        MyEntityA objA = new MyEntityA();
        MyEntityB objB = new MyEntityB();
        objB.observe(objA);

        new MyThreadA(objA).start();
        new MyThreadB(objB).start();

    }
}

```

AVOIDING DEADLOCKS

- A program that never acquires more than one lock a time cannot have lock-ordering deadlock
- if we must acquire multiple locks, lock ordering must be part of the design
 - minimizing the number of potential locking interaction
 - document a lock-ordering protocol for locks that may be acquired together
- Alternative technique: *timed locks*
 - detecting and recovering from deadlocks using `tryLock` feature of `Lock` classes instead of intrinsic lock

OTHER LIVENESS HAZARD

- **Starvation**
 - typically manifested when using *priorities*
 - basic thread support for priorities in Java thread is “deprecated”
 - platform dependent
 - liveness problems
- **Poor responsiveness**
 - e.g. executing long-term tasks in GUI thread
 - can also be caused by *poor lock management*
 - if a thread holds a lock for long time - for instance while iterating on a large collection and performing substantial work on each element - other threads that need to access that collection may have to wait long time
- **Livelock**
 - when threads cannot make progress because they keep retrying an operation that will always fail
 - solution: introducing some *randomness* into the retry mechanism
 - breaking the synchronization that causes the live-lock

BIBLIOGRAPHY

- *“Java Concurrency in Practice”*, Brian Goetz, Addison Wesley